


# Large Scale Data Management

José Orlando Pereira

Departamento de Informática  
Universidade do Minho



# Big Data

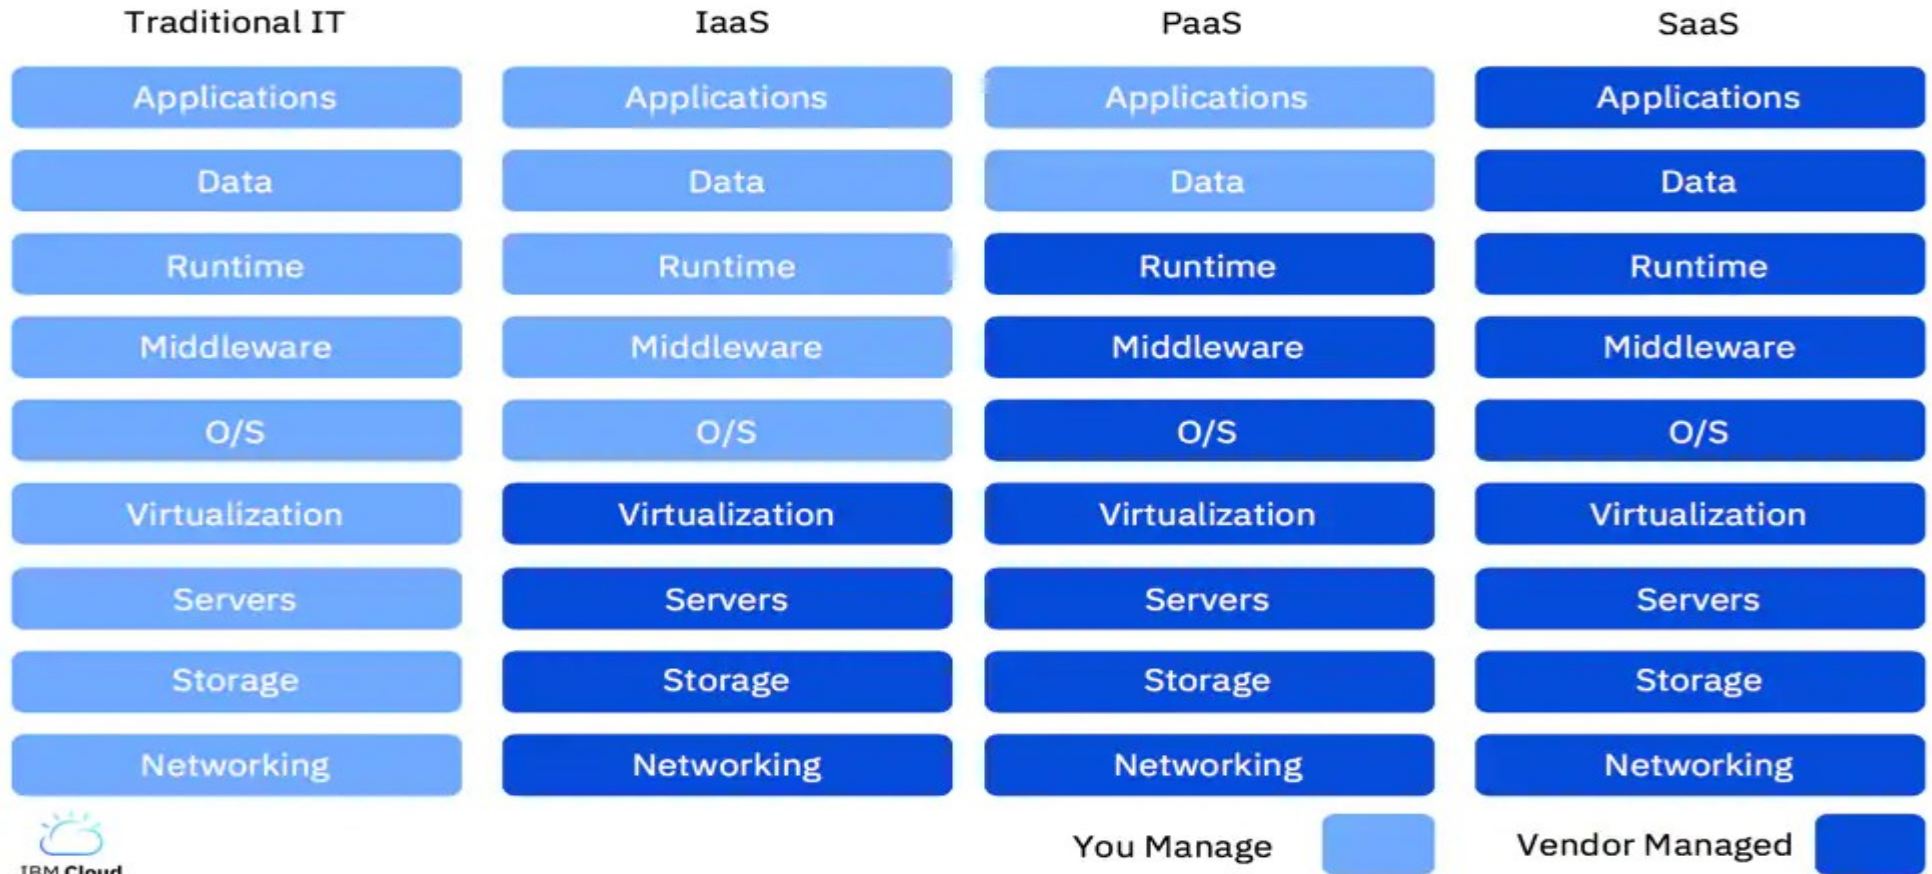


Sources: McKinsey Global Institute, Twitter, Cisco, Gartner, EMC, SAS, IBM, MEPTec, QAS

<https://www.ibmbigdatahub.com/infographic/four-vs-big-data>

IBM

# Cloud computing



<https://www.ibm.com/cloud/learn/iaas>

# Contents

- Distributed processing (Map-Reduce)
- Scalable storage (HDFS and HBase)
- Data flow and streaming (Spark)

# Main references

- M. Tamer Özsu, Patrick Valduriez. **Principles of Distributed Database Systems** (3rd Edition). Springer.
  - Chapter 18
- Peter Bailis, Joseph M. Hellerstein, Michael Stonebraker. **Readings in Database Systems** (5th Edition)
  - Chapter 5: <http://www.redbook.io/ch5-dataflow.html>
    - <https://static.googleusercontent.com/media/research.google.com/en//archive/mapreduce-osdi04.pdf>
    - <https://storageconference.us/2010/Papers/MSST/Shvachko.pdf>
    - <https://storage.googleapis.com/pub-tools-public-publication-data/pdf/68a74a85e1662fe02ff3967497f31fda7f32225c.pdf>
    - [http://people.csail.mit.edu/matei/papers/2012/nsdi\\_spark.pdf](http://people.csail.mit.edu/matei/papers/2012/nsdi_spark.pdf)

# Additional references

- Tom White. **Hadoop: The Definitive Guide** (4th Edition). O'Reilly.
- Lars George. **HBase: The Definitive Guide**. O'Reilly.
- Holden Karau, Andy Konwinski, Patrick Wendell & Matei Zaharia. **Spark: Lightning-fast data analysis**. O'Reilly.
- Software manuals

# Grading

- Group projects
  - First project (MapReduce)
  - Second project (Spark)
  - 50% weight
- Written exam
  - 50% weight
  - Minimum: 8 / 20





# Maven Build

- Automatic dependency downloading and packaging
- Set Java version:

```
<properties>  
  <maven.compiler.source>8</maven.compiler.source>  
  <maven.compiler.target>8</maven.compiler.target>  
</properties>
```

- Packaging:

```
<plugin>  
  <groupId>org.apache.maven.plugins</groupId>  
  <artifactId>maven-jar-plugin</artifactId>  
  <configuration>  
    <archive>  
      <manifest>  
        <addClasspath>true</addClasspath>  
        <classpathPrefix>libs/</classpathPrefix>  
        <mainClass>  
          packagename.MainClass  
        </mainClass>  
      </manifest>  
    </archive>  
  </configuration>  
</plugin>
```


# Maven Build

- Sample dependency:

```
<dependency>
  <groupId>org.apache.commons</groupId>
  <artifactId>commons-compress</artifactId>
  <version>1.20</version>
</dependency>
```

- Dependency packaging:

```
<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-shade-plugin</artifactId>
  <version>3.2.2</version>
  <executions>
    <execution>
      <phase>package</phase>
      <goals>
        <goal>shade</goal>
      </goals>
      <configuration>
        <transformers>
          <transformer implementation=
"org.apache.maven.plugins.shade.resource.Services
ResourceTransformer"/>
        </transformers>
      </configuration>
    </execution>
  </executions>
</plugin>
```



docker

# Docker Setup

- Installation guide:

<https://docs.docker.com/get-docker/>

(do not skip “Post-installation steps for Linux”!!!)

- Testing:

```
$ docker ps
```

```
$ docker pull hello-world
```

```
$ docker run hello-world
```

```
$ docker ps -a
```

```
$ docker rm containername
```

# Docker Basics

- Running a container:

```
$ docker run -it ubuntu
```

```
$ docker run -it --name containername ubuntu
```

- Entering into an existing container

```
$ docker ps
```

```
$ docker exec -it containername bash
```

- Stopping and cleaning up:

```
$ docker kill containername
```

```
$ docker ps -a
```

```
$ docker rm containername
```

# Packaging a Java application

- Sample Dockerfile:

```
FROM openjdk:8
COPY target/jarname.jar /
ENTRYPOINT ["java", "-jar", "/jarname.jar"]
```

- Build with:

```
$ docker build -t imagename .
```

```
$ docker images
```

- Run with:

```
$ docker run -it --name conthname imagename arg1 arg2 ...
```

```
$ docker rm conthname
```

```
$ docker rmi imagename
```

# Accessing host files

- The container and the host have separate file-systems
  - Map a folder from the host into the container
  - Reference files in that folder

```
$ docker run -it \
```

```
    -v /hostfolder:/containerfolder
```

```
    imagename /containerfolder/filename
```